

chronic kidney disease

AUTHOR

Kenzie_lee

chronic kidney disease dataset analysis

data preparation

```
library(caret)
library(ggplot2)
library(corrplot)
library(dplyr)
library(reshape2)
library(glmnet)
```

```
file_path <- '/Users/apple/Desktop/me/WCM/semester/fall1/biosta1/final/kidney_disease.csv'
df <- read.csv(file_path)
#str(df)
head(df)
```

	Patient.ID	Age.of.the.patient	Blood.Pressure	Specific.Gravity..urine.
1	id,k	age	bp	sg
2	0	48	80	1.02
3	1	7	50	1.02
4	2	62	80	1.01
5	3	48	70	1.005
6	4	51	80	1.01
	Albumin.in.urine	Sugar.in.urine	Red.Blood.Cells.in.urine	Pus.Cells.in.urine
1	al	su	rbc	pc
2	1	0		normal
3	4	0		normal
4	2	3	normal	normal
5	4	0	normal	abnormal
6	2	0	normal	normal
	Pus.Cell.Clumps	Bacteria.in.urine	Blood.Glucose.Random	Blood.Urea
1	pcc	ba	bgr	bu
2	notpresent	notpresent	121	36
3	notpresent	notpresent		18
4	notpresent	notpresent	423	53
5	present	notpresent	117	56
6	notpresent	notpresent	106	26
	Serum.Creatinine	Sodium.in.blood	Potassium.in.blood	Hemoglobin.level
1	sc	sod	pot	hemo
2	1.2			15.4
3	0.8			11.3
4	1.8			9.6
5	3.8	111	2.5	11.2
6	1.4			11.6
	Packed.Cell.Volume..Hematocrit.	White.Blood.Cell.Count	Red.Blood.Cell.Count	

1		pcv	wc	rc
2		44	7800	5.2
3		38	6000	
4		31	7500	
5		32	6700	3.9
6		35	7300	4.6
Hypertension Diabetes.Mellitus Coronary.Artery.Disease Appetite				
1	htn	dm	cad	appet
2	yes	yes	no	good
3	no	no	no	good
4	no	yes	no	poor
5	yes	no	no	poor
6	no	no	no	good
Pedal.Edema..swelling.of.legs. Anemia Diagnosis.label				
1	pe	ane	classification	
2	no	no	ckd	
3	no	no	ckd	
4	no	yes	ckd	
5	yes	yes	ckd	
6	no	no	ckd	

First check duplicated rows.

```
#check duplicated rows
sum(duplicated(df))
```

[1] 0

```
#df[duplicated(df), ]
```

Delete extra columns, such as the ID column. Classify the numerical and categorical data.

```
df <- subset(df, select = -Patient.ID)#please run this for one first time
#df <- df[, !names(df) %in% c("id")]
#str(df)
#colnames(df)

#name the numerical and categorical data
num_cols <- c(
  "Age.of.the.patient",
  "Blood.Pressure",
  "Specific.Gravity..urine.",
  "Blood.Glucose.Random",
  "Blood.Urea",
  "Serum.Creatinine",
  "Sodium.in.blood",
  "Potassium.in.blood",
  "Hemoglobin.level",
  "Packed.Cell.Volume..Hematocrit.",
  "White.Blood.Cell.Count",
  "Red.Blood.Cell.Count"
)

df[num_cols] <- lapply(df[num_cols], function(x) as.numeric(x))
```

```
str(df)
```

```
'data.frame': 401 obs. of 25 variables:
 $ Age.of.the.patient      : num NA 48 7 62 48 51 60 68 24 52 ...
 $ Blood.Pressure          : num NA 80 50 80 70 80 90 70 NA 100 ...
 $ Specific.Gravity..urine: num NA 1.02 1.02 1.01 1 ...
 $ Albumin.in.urine        : chr "al" "1" "4" "2" ...
 $ Sugar.in.urine          : chr "su" "0" "0" "3" ...
 $ Red.Blood.Cells.in.urine: chr "rbc" "" "" "normal" ...
 $ Pus.Cells.in.urine       : chr "pc" "normal" "normal" "normal" ...
 $ Pus.Cell.Clumps          : chr "pcc" "notpresent" "notpresent" "notpresent"
 ...
 $ Bacteria.in.urine        : chr "ba" "notpresent" "notpresent" "notpresent"
 ...
 $ Blood.Glucose.Random     : num NA 121 NA 423 117 106 74 100 410 138 ...
 $ Blood.Urea               : num NA 36 18 53 56 26 25 54 31 60 ...
 $ Serum.Creatinine         : num NA 1.2 0.8 1.8 3.8 1.4 1.1 24 1.1 1.9 ...
 $ Sodium.in.blood          : num NA NA NA NA 111 NA 142 104 NA NA ...
 $ Potassium.in.blood        : num NA NA NA NA 2.5 NA 3.2 4 NA NA ...
 $ Hemoglobin.level          : num NA 15.4 11.3 9.6 11.2 11.6 12.2 12.4 12.4
10.8 ...
 $ Packed.Cell.Volume..Hematocrit.: num NA 44 38 31 32 35 39 36 44 33 ...
 $ White.Blood.Cell.Count    : num NA 7800 6000 7500 6700 7300 7800 NA 6900 9600
...
 $ Red.Blood.Cell.Count      : num NA 5.2 NA NA 3.9 4.6 4.4 NA 5 4 ...
 $ Hypertension              : chr "htn" "yes" "no" "no" ...
 $ Diabetes.Mellitus         : chr "dm" "yes" "no" "yes" ...
 $ Coronary.Artery.Disease   : chr "cad" "no" "no" "no" ...
 $ Appetite                  : chr "appet" "good" "good" "poor" ...
 $ Pedal.Edema..swelling.of.legs.: chr "pe" "no" "no" "no" ...
 $ Anemia                    : chr "ane" "no" "no" "yes" ...
 $ Diagnosis.label            : chr "classification" "ckd" "ckd" "ckd" ...
```

```
cat_cols <- names(df)[sapply(df, is.character) | sapply(df, is.factor)]
#num_col <- names(df)[!(sapply(df, is.character) | sapply(df, is.factor))]
#
cat_cols
```

```
[1] "Albumin.in.urine"           "Sugar.in.urine"
[3] "Red.Blood.Cells.in.urine"   "Pus.Cells.in.urine"
[5] "Pus.Cell.Clumps"           "Bacteria.in.urine"
[7] "Hypertension"               "Diabetes.Mellitus"
[9] "Coronary.Artery.Disease"    "Appetite"
[11] "Pedal.Edema..swelling.of.legs." "Anemia"
[13] "Diagnosis.label"
```

```
#num_col
```

```
#summary(df)
```

Deal with "tab".

```

df <- df[-1,]
#delete the first row: df <- df[-1, ]
#delete the first two rows: df <- df[-c(1,2), ]
#delete the i th row: df <- df[-i, ]
#remain 2 to all rows df <- df[2:nrow(df), ]
#head(df)

df$`Diabetes.Mellitus` <- gsub("\\tno", "no", df$`Diabetes.Mellitus`)
df$`Diabetes.Mellitus` <- gsub("\\tyes", "yes", df$`Diabetes.Mellitus`)

df$`Coronary.Artery.Disease` <- gsub("\\tno", "no", df$`Coronary.Artery.Disease`)
df$`Coronary.Artery.Disease` <- gsub("\\tyes", "yes", df$`Coronary.Artery.Disease`)

df$Diagnosis.label <- gsub("ckd\\t", "ckd", df$Diagnosis.label )
#summary(df)

```

Deal with NAs.

```
colSums(is.na(df))
```

Age.of.the.patient	Blood.Pressure
9	12
Specific.Gravity..urine.	Albumin.in.urine
47	0
Sugar.in.urine	Red.Blood.Cells.in.urine
0	0
Pus.Cells.in.urine	Pus.Cell.Clumps
0	0
Bacteria.in.urine	Blood.Glucose.Random
0	44
Blood.Urea	Serum.Creatinine
19	17
Sodium.in.blood	Potassium.in.blood
87	88
Hemoglobin.level	Packed.Cell.Volume..Hematocrit.
52	71
White.Blood.Cell.Count	Red.Blood.Cell.Count
106	131
Hypertension	Diabetes.Mellitus
0	0
Coronary.Artery.Disease	Appetite
0	0
Pedal.Edema..swelling.of.legs.	Anemia
0	0
Diagnosis.label	
0	

Missing value imputation methods:

- For numerical variables (e.g., age, blood pressure, blood sugar): Randomly select some values from the known non-missing values of the column and replace the missing values. • If the mean or median is used for imputation, all missing values will be replaced with the same number, which will reduce the variance of the variable and cause the data to be "oversmoothed". Random

imputation allows the imputed values to maintain the true distribution characteristics of the data (e.g., the probability that people with high blood sugar will take higher values and people with low blood sugar will take lower values is still preserved). After this processing, the model training will not lose information diversity due to a large number of identical values.

2. For categorical variables: We replace missing values with the value that appears most frequently (mode). These variables have no continuous relationship (they cannot be averaged or randomly selected), and mode imputation is the most natural choice because it preserves the main distribution of the data; it does not introduce invalid categories; and it has the least impact on model bias.

```
#random -- NA--numeric
Random_value_Imputation <- function(x, seed=42) {
  set.seed(seed)
  na_index <- is.na(x)
  x[na_index] <- sample(x[!na_index], sum(na_index), replace=TRUE)
  return(x)
}# "Albumin.in.urine"           "Sugar.in.urine"

#mode -- NA--cate
impute_mode <- function(x, seed=42) {
  set.seed(seed)
  mode_val <- names(sort(table(x), decreasing = TRUE))[1]
  x[is.na(x)] <- mode_val
  return(x)
}

#num_cols
#cat_cols
for (col in num_cols) {
  df[[col]] <- Random_value_Imputation(df[[col]], seed=203)
}
for (col in cat_cols) {
  df[[col]] <- impute_mode(df[[col]], seed=203)
}
```

Draw a heatmap.

Drawing a correlation plot is essential: If a high correlation is found ($|r| > 0.8$), it is recommended to remove one of the variables; If there are many variables, run vif() first and then delete them; Variable selection is best done after cleaning and before modeling; If you want to predict (rather than explain), you can use penalized regression for automatic selection.

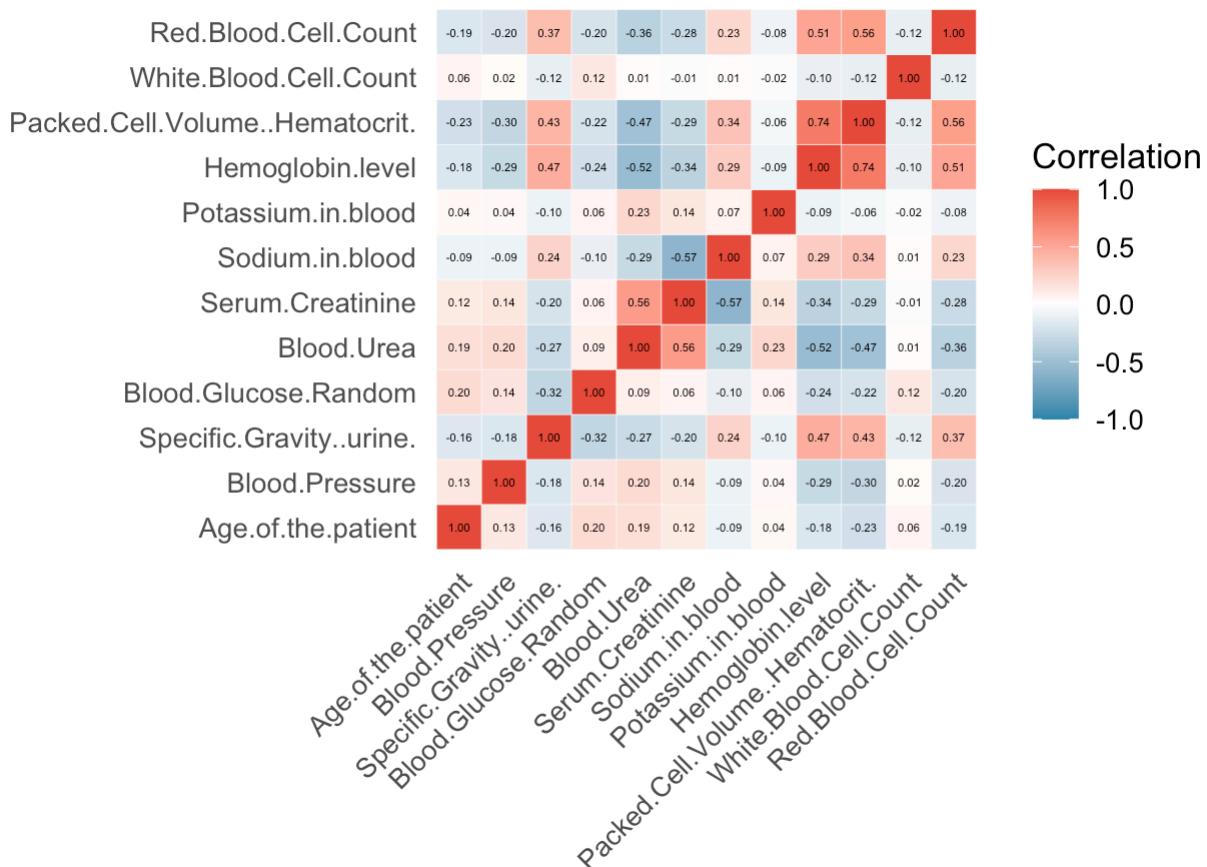
```
num_data <- df[, num_cols]

#make sure it's all numeric
num_data <- mutate_all(num_data, function(x) as.numeric(as.character(x)))

#complete.obs omits NA
corr_matrix <- cor(num_data, use = "complete.obs")
#turn it into matrix
corr_melt <- melt(corr_matrix)
```

```
ggplot(corr_melt, aes(x = Var1, y = Var2, fill = value)) +
  geom_tile(color = "white") +
  scale_fill_gradient2(low = "#2E86AB", mid = "white", high = "#E74C3C",
                       midpoint = 0, limit = c(-1, 1), name = "Correlation") +
  theme_minimal(base_size = 13) +
  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust = 1),
        panel.grid = element_blank(),
        axis.title = element_blank()) +
  coord_fixed() +
  geom_text(aes(label = sprintf("%.2f", value)), size = 1.5, color = "black") +
  labs(title = "Correlation Heatmap of Numeric Variables")
```

Correlation Heatmap of Numeric Variables



Perform label encoding for categorical variables

```
summary(df)
```

Age.of.the.patient	Blood.Pressure	Specific.Gravity..urine.
Min. : 2.00	Min. : 50.00	Min. : 1.005
1st Qu.: 42.00	1st Qu.: 70.00	1st Qu.: 1.010
Median : 54.00	Median : 80.00	Median : 1.020
Mean : 51.56	Mean : 76.53	Mean : 1.017
3rd Qu.: 64.25	3rd Qu.: 80.00	3rd Qu.: 1.020
Max. : 90.00	Max. : 180.00	Max. : 1.025
Albumin.in.urine	Sugar.in.urine	Red.Blood.Cells.in.urine
Length:400	Length:400	Length:400
Class : character	Class : character	Class : character

```
Mode :character Mode :character Mode :character
```

Pus.Cells.in.urine	Pus.Cell.Clumps	Bacteria.in.urine	Blood.Glucose.Random
Length:400	Length:400	Length:400	Min. : 22.0
Class :character	Class :character	Class :character	1st Qu.:100.0
Mode :character	Mode :character	Mode :character	Median :122.0
			Mean :149.9
			3rd Qu.:169.0
			Max. :490.0
Blood.Urea	Serum.Creatinine	Sodium.in.blood	Potassium.in.blood
Min. : 1.50	Min. : 0.400	Min. : 4.5	Min. : 2.500
1st Qu.: 27.00	1st Qu.: 0.900	1st Qu.:135.0	1st Qu.: 3.800
Median : 44.00	Median : 1.300	Median :138.0	Median : 4.400
Mean : 58.65	Mean : 3.100	Mean :137.6	Mean : 4.774
3rd Qu.: 68.00	3rd Qu.: 2.825	3rd Qu.:142.0	3rd Qu.: 4.900
Max. :391.00	Max. :76.000	Max. :163.0	Max. :47.000
Hemoglobin.level	Packed.Cell.Volume..Hematocrit.	White.Blood.Cell.Count	
Min. : 3.10	Min. : 9.00	Min. : 2200	
1st Qu.:10.30	1st Qu.:32.75	1st Qu.: 6400	
Median :12.60	Median :40.00	Median : 8000	
Mean :12.47	Mean :38.92	Mean : 8357	
3rd Qu.:15.00	3rd Qu.:45.00	3rd Qu.: 9800	
Max. :17.80	Max. :54.00	Max. :26400	
Red.Blood.Cell.Count	Hypertension	Diabetes.Mellitus	
Min. :2.100	Length:400	Length:400	
1st Qu.:3.900	Class :character	Class :character	
Median :4.700	Mode :character	Mode :character	
Mean :4.709			
3rd Qu.:5.400			
Max. :8.000			
Coronary.Artery.Disease	Appetite	Pedal.Edema..swelling.of.legs.	
Length:400	Length:400	Length:400	
Class :character	Class :character	Class :character	
Mode :character	Mode :character	Mode :character	

Anemia	Diagnosis.label
Length:400	Length:400
Class :character	Class :character
Mode :character	Mode :character

```
#colSums(is.na(df))
head(df)
```

Age.of.the.patient	Blood.Pressure	Specific.Gravity..urine.	Albumin.in.urine	
2	48	80	1.020	1
3	7	50	1.020	4

4	62	80	1.010	2
5	48	70	1.005	4
6	51	80	1.010	2
7	60	90	1.015	3
Sugar.in.urine Red.Blood.Cells.in.urine Pus.Cells.in.urine Pus.Cell.Clumps				
2	0	normal	notpresent	
3	0	normal	notpresent	
4	3	normal	normal	notpresent
5	0	normal	abnormal	present
6	0	normal	normal	notpresent
7	0			notpresent
Bacteria.in.urine Blood.Glucose.Random Blood.Urea Serum.Creatinine				
2	notpresent	121	36	1.2
3	notpresent	230	18	0.8
4	notpresent	423	53	1.8
5	notpresent	117	56	3.8
6	notpresent	106	26	1.4
7	notpresent	74	25	1.1
Sodium.in.blood Potassium.in.blood Hemoglobin.level				
2	144	5.0	15.4	
3	140	3.5	11.3	
4	140	4.0	9.6	
5	111	2.5	11.2	
6	147	4.1	11.6	
7	142	3.2	12.2	
Packed.Cell.Volume..Hematocrit. White.Blood.Cell.Count Red.Blood.Cell.Count				
2		44	7800	5.2
3		38	6000	5.8
4		31	7500	3.7
5		32	6700	3.9
6		35	7300	4.6
7		39	7800	4.4
Hypertension Diabetes.Mellitus Coronary.Artery.Disease Appetite				
2	yes	yes	no	good
3	no	no	no	good
4	no	yes	no	poor
5	yes	no	no	poor
6	no	no	no	good
7	yes	yes	no	good
Pedal.Edema..swelling.of.legs. Anemia Diagnosis.label				
2		no no	ckd	
3		no no	ckd	
4		no yes	ckd	
5		yes yes	ckd	
6		no no	ckd	
7		yes no	ckd	

```
#df<-subset(df,select = -Patient.ID)
cat_cols <- setdiff(cat_cols, "Patient.ID")

# Perform label encoding for categorical variables
df$Hypertension <- as.integer(factor(df$Hypertension))
df$Diabetes.Mellitus <- as.integer(factor(df$Diabetes.Mellitus))
```

```

df$Coronary.Artery.Disease <- as.integer(factor(df$Coronary.Artery.Disease))
df$Appetite <- as.integer(factor(df$Appetite))
df$Pedal.Edema..swelling.of.legs. <- as.integer(factor(df$Pedal.Edema..swelling.of.leg
df$Anemia <- as.integer(factor(df$Anemia))
df$Red.Blood.Cells.in.urine <- as.integer(factor(df$Red.Blood.Cells.in.urine))
df$Albumin.in.urine <- as.integer(factor(df$Albumin.in.urine))
df$Sugar.in.urine <- as.integer(factor(df$Sugar.in.urine))
df$Pus.Cells.in.urine <- as.integer(factor(df$Pus.Cells.in.urine))
df$Pus.Cell.Clumps <- as.integer(factor(df$Pus.Cell.Clumps))
df$Bacteria.in.urine <- as.integer(factor(df$Bacteria.in.urine))
df$ Diagnosis.label <- as.integer(factor(df$ Diagnosis.label))
cat_cols

```

```

[1] "Albumin.in.urine"           "Sugar.in.urine"
[3] "Red.Blood.Cells.in.urine"   "Pus.Cells.in.urine"
[5] "Pus.Cell.Clumps"          "Bacteria.in.urine"
[7] "Hypertension"              "Diabetes.Mellitus"
[9] "Coronary.Artery.Disease"    "Appetite"
[11] "Pedal.Edema..swelling.of.legs." "Anemia"
[13] "Diagnosis.label"

```

```

#str(df)

sapply(df, is.numeric)

```

Age.of.the.patient	Blood.Pressure
TRUE	TRUE
Specific.Gravity..urine.	Albumin.in.urine
TRUE	TRUE
Sugar.in.urine	Red.Blood.Cells.in.urine
TRUE	TRUE
Pus.Cells.in.urine	Pus.Cell.Clumps
TRUE	TRUE
Bacteria.in.urine	Blood.Glucose.Random
TRUE	TRUE
Blood.Urea	Serum.Creatinine
TRUE	TRUE
Sodium.in.blood	Potassium.in.blood
TRUE	TRUE
Hemoglobin.level	Packed.Cell.Volume..Hematocrit.
TRUE	TRUE
White.Blood.Cell.Count	Red.Blood.Cell.Count
TRUE	TRUE
Hypertension	Diabetes.Mellitus
TRUE	TRUE
Coronary.Artery.Disease	Appetite
TRUE	TRUE
Pedal.Edema..swelling.of.legs.	Anemia
TRUE	TRUE
Diagnosis.label	
TRUE	

```

index <- createDataPartition(df$ Diagnosis.label, p = 0.8, list = FALSE)
train_data <- df[index, ]
test_data <- df[-index, ]

#Convert tags to binary; machine learning models typically require binary tags.
train_data$ Diagnosis.label <- ifelse(train_data$ Diagnosis.label == 2, 1, 0)
test_data$ Diagnosis.label <- ifelse(test_data$ Diagnosis.label == 2, 1, 0)

#split X / y
X_train <- as.matrix(train_data[, -which(names(train_data) == "Diagnosis.label")])
X_test <- as.matrix(test_data[, -which(names(test_data) == "Diagnosis.label")])
y_train <- train_data$ Diagnosis.label
y_test <- test_data$ Diagnosis.label

#summary(train_data)

```

GLM Modeling

```

log_model <- glm( Diagnosis.label ~ ., data = train_data, family = binomial)
y_pred_prob <- predict(log_model, test_data, type = "response")
y_pred <- ifelse(y_pred_prob > 0.5, 1, 0)

acc <- mean(y_pred == test_data$ Diagnosis.label)

print(acc)

```

[1] 1

CKD data is typical: indicators like urine protein (albumin-in-urine), urine glucose, and pus cells in urine, once they reach a certain level, almost always indicate CKD; while normal levels almost never indicate CKD. This is the worst-case scenario for logistic regression: one or two variables that can almost separate 0 and 1.

The simplest approach to this type of problem is to use `glmnet`, not the bare `glm()`. This is because the penalty term prevents the coefficients from going to infinity.

```

#make matrix
x_train <- model.matrix(Diagnosis.label ~ ., data = train_data)[, -1]
y_train <- train_data$Diagnosis.label

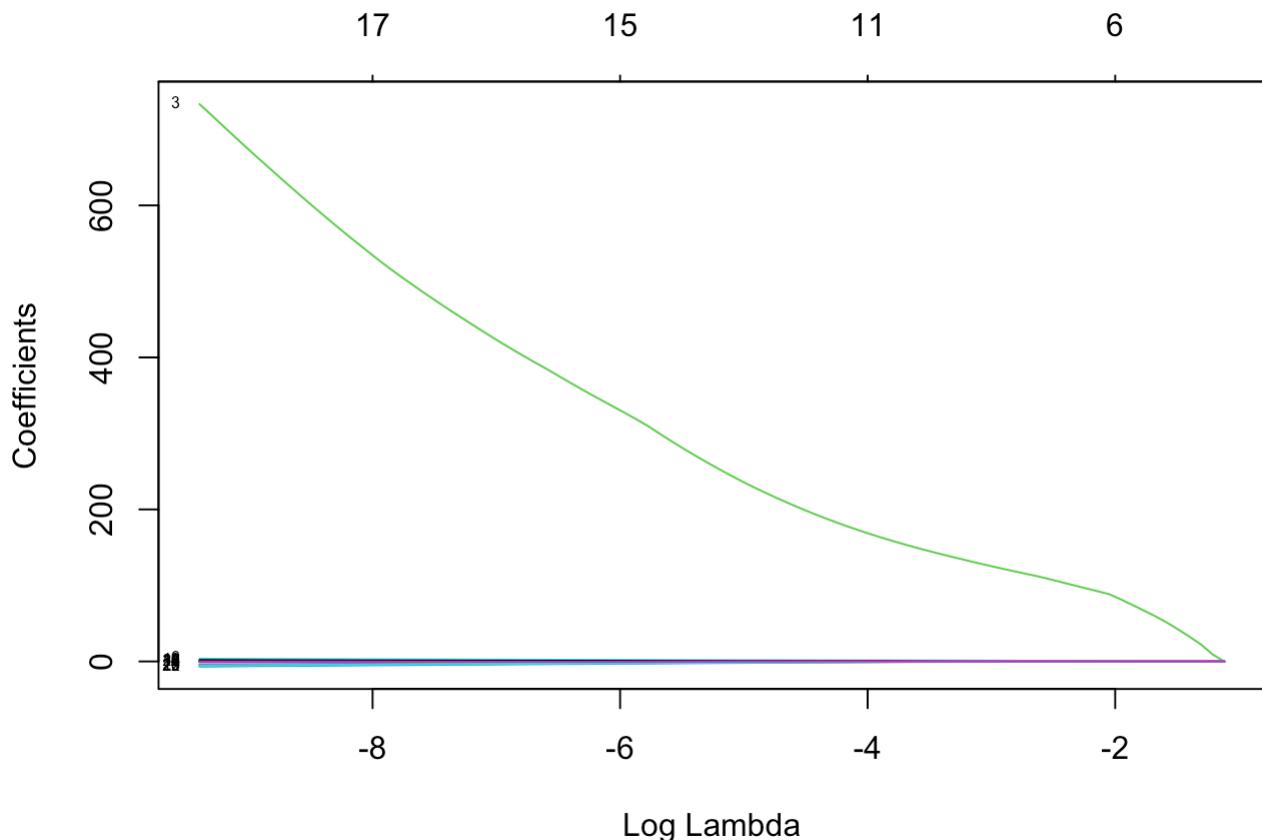
fit <- glmnet(x_train, y_train, family = "binomial")

#predict
x_test <- model.matrix(Diagnosis.label ~ ., data = test_data)[, -1]
y_pred_prob <- predict(fit, newx = x_test, s = 0.01, type = "response")
y_pred <- ifelse(y_pred_prob > 0.5, 1, 0)
mean(y_pred == test_data$Diagnosis.label)

```

[1] 1

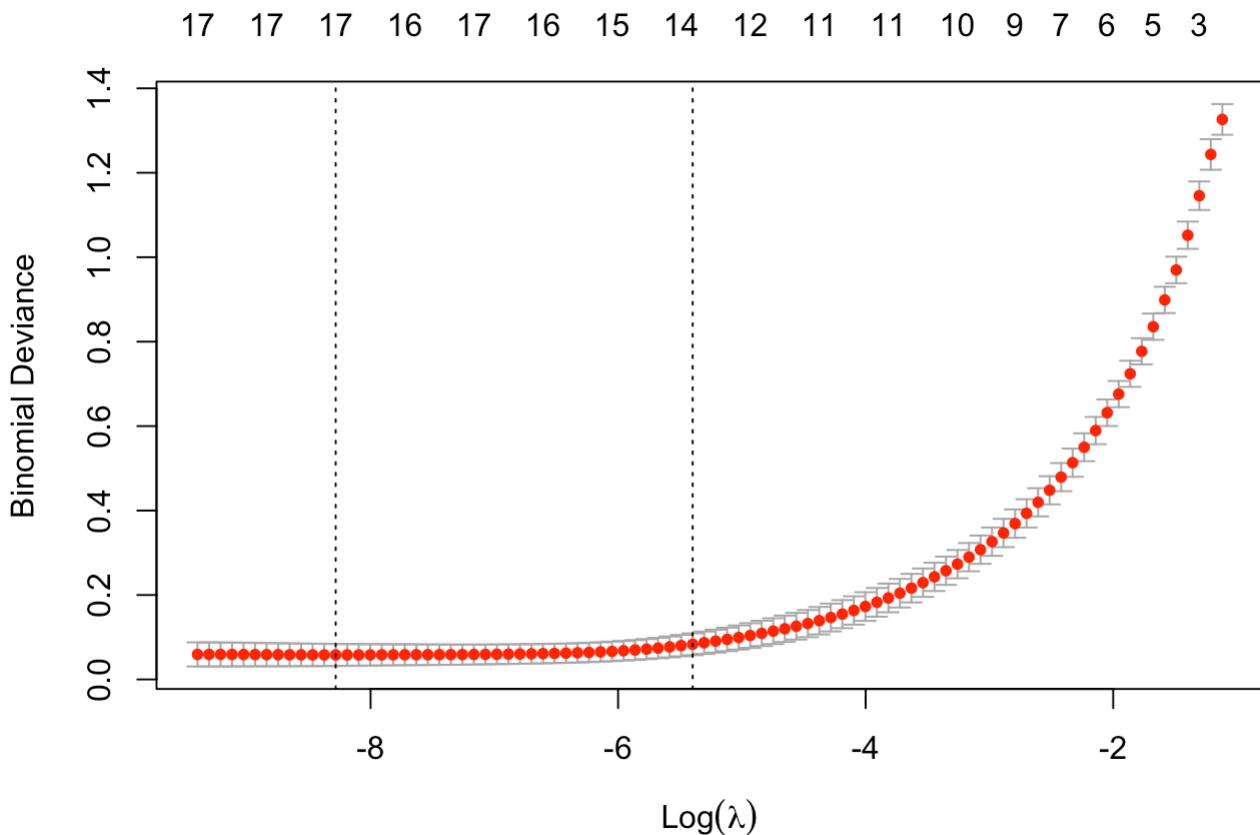
```
plot(fit, xvar = "lambda", label = TRUE)
```



About glm choosing lambda.

```
set.seed(50)
cvfit <- cv.glmnet(x_train, y_train, family = "binomial", alpha = 1)

plot(cvfit)
```



```
cvfit$lambda.min #best λ
```

```
[1] 0.0002530825
```

```
cvfit$lambda.1se #stable λ
```

```
[1] 0.004526759
```

left → lambda.min (λ with the smallest error) right → lambda.1se (easiest and stable λ)

Linear regression

```
dat <- df
cat_cols
```

```
[1] "Albumin.in.urine"           "Sugar.in.urine"
[3] "Red.Blood.Cells.in.urine"    "Pus.Cells.in.urine"
[5] "Pus.Cell.Clumps"            "Bacteria.in.urine"
[7] "Hypertension"                "Diabetes.Mellitus"
[9] "Coronary.Artery.Disease"     "Appetite"
[11] "Pedal.Edema..swelling.of.legs." "Anemia"
[13] "Diagnosis.label"
```

```
#split(num_cols,f = ",",sep = "+")
#find names
```

```
paste(cat_cols, collapse = "+")
```

```
[1]
"Albumin.in.urine+Sugar.in.urine+Red.Blood.Cells.in.urine+Pus.Cells.in.urine+Pus.Cell.
Clumps+Bacteria.in.urine+Hypertension+Diabetes.Mellitus+Coronary.Artery.Disease+Appeti
te+Pedal.Edema..swelling.of.legs.+Anemia+Diagnosis.label"
```

```
#independent variables
lm_fit <- lm(Diagnosis.label ~ Hemoglobin.level+Albumin.in.urine+Sugar.in.urine+Red.B
data = dat)

summary(lm_fit)
```

Call:

```
lm(formula = Diagnosis.label ~ Hemoglobin.level + Albumin.in.urine +
  Sugar.in.urine + Red.Blood.Cells.in.urine + Pus.Cells.in.urine +
  Pus.Cell.Clumps + Bacteria.in.urine + Hypertension + Diabetes.Mellitus +
  Coronary.Artery.Disease + Appetite + Pedal.Edema..swelling.of.legs. +
  Anemia + Serum.Creatinine, data = dat)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.72222	-0.14429	0.02754	0.13604	0.70115

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.161819	0.234891	4.946	1.13e-06 ***
Hemoglobin.level	0.045430	0.005536	8.206	3.48e-15 ***
Albumin.in.urine	-0.100269	0.010891	-9.207	< 2e-16 ***
Sugar.in.urine	-0.021568	0.011912	-1.811	0.07098 .
Red.Blood.Cells.in.urine	0.245296	0.014398	17.037	< 2e-16 ***
Pus.Cells.in.urine	0.037811	0.018470	2.047	0.04132 *
Pus.Cell.Clumps	-0.049330	0.041120	-1.200	0.23101
Bacteria.in.urine	-0.042472	0.051083	-0.831	0.40625
Hypertension	-0.089461	0.033887	-2.640	0.00863 **
Diabetes.Mellitus	-0.071411	0.031744	-2.250	0.02504 *
Coronary.Artery.Disease	0.069109	0.043929	1.573	0.11650
Appetite	-0.062216	0.032872	-1.893	0.05915 .
Pedal.Edema..swelling.of.legs.	-0.013195	0.034795	-0.379	0.70474
Anemia	0.016579	0.037287	0.445	0.65684
Serum.Creatinine	-0.004483	0.002192	-2.045	0.04156 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2278 on 385 degrees of freedom

Multiple R-squared: 0.787, Adjusted R-squared: 0.7792

F-statistic: 101.6 on 14 and 385 DF, p-value: < 2.2e-16

```
car::vif(lm_fit)#whether there is a multicollinearity
```

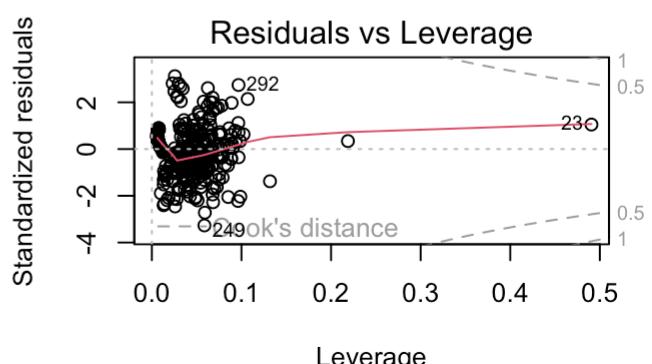
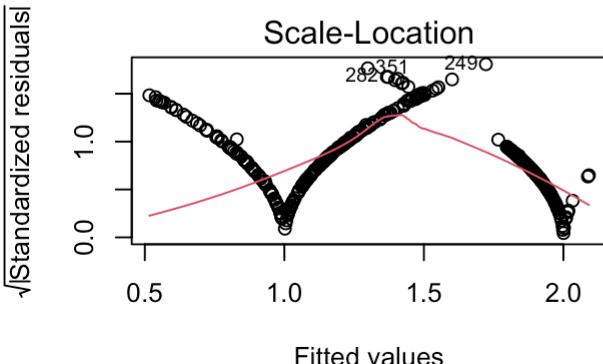
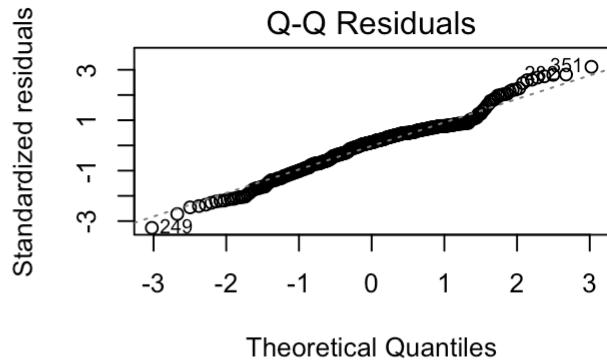
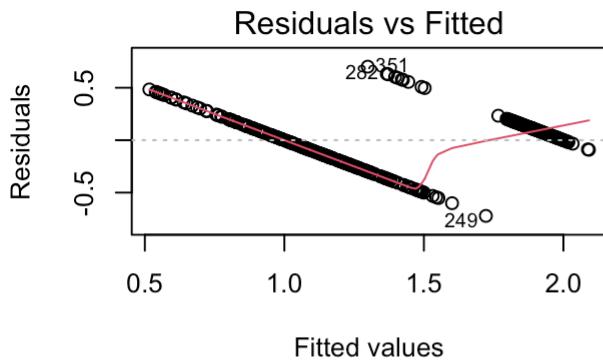
Hemoglobin.level	2.088912	Albumin.in.urine	1.855492
Sugar.in.urine	1.403951	Red.Blood.Cells.in.urine	1.386586
Pus.Cells.in.urine	1.511873	Pus.Cell.Clumps	1.381603
Bacteria.in.urine	1.267104	Hypertension	2.134647
Diabetes.Mellitus	1.983162	Coronary.Artery.Disease	1.243940
Appetite	1.387175	Pedal.Edema..swelling.of.legs.	1.468771
Anemia	1.401583	Serum.Creatinine	1.193711

```
#VIF > 10: Severe collinearity;
```

```
#5 < VIF ≤ 10: Moderate;
```

```
#VIF < 5: Acceptable.
```

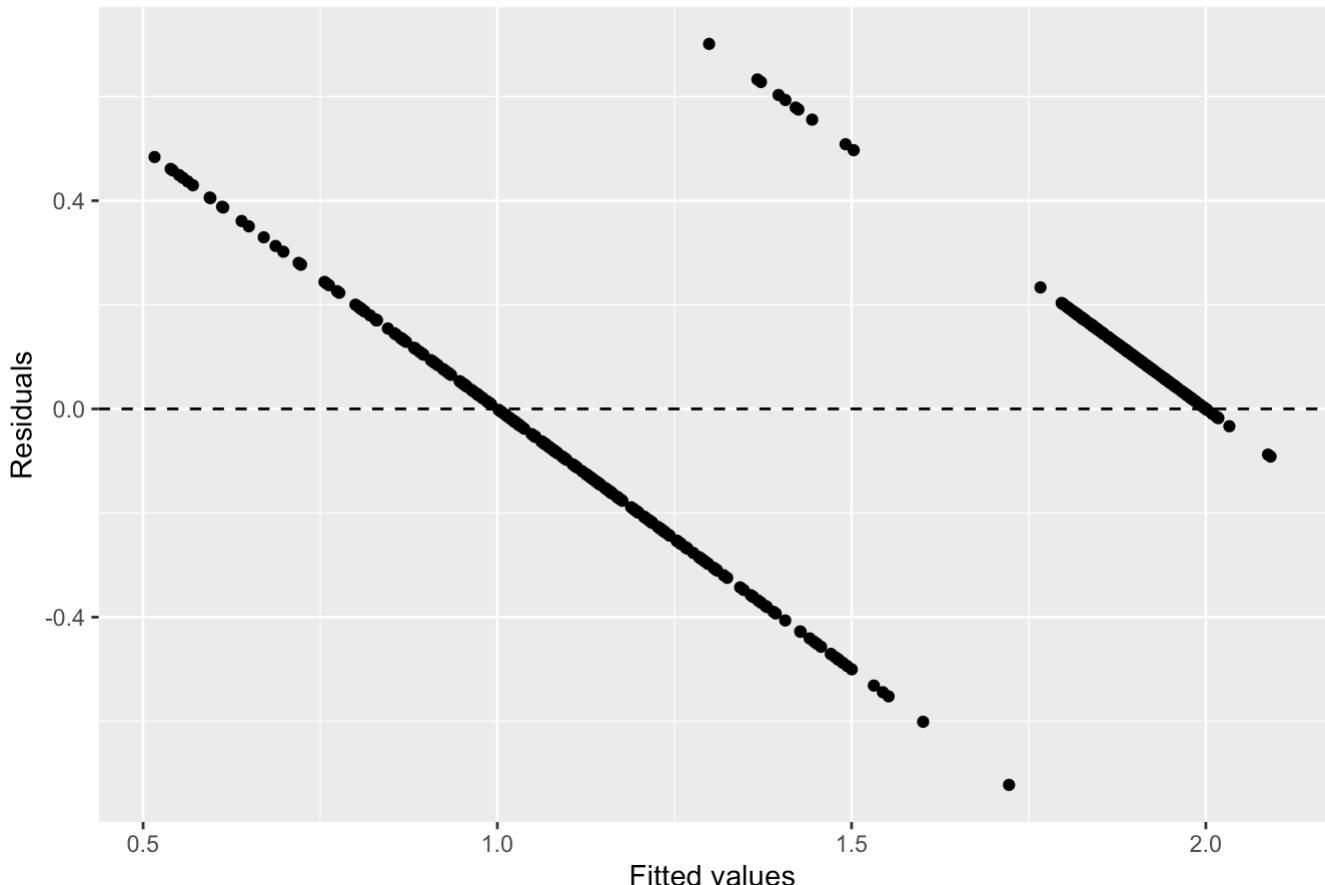
```
par(mfrow = c(2, 2))
plot(lm_fit)
```



```
par(mfrow = c(1, 1))
```

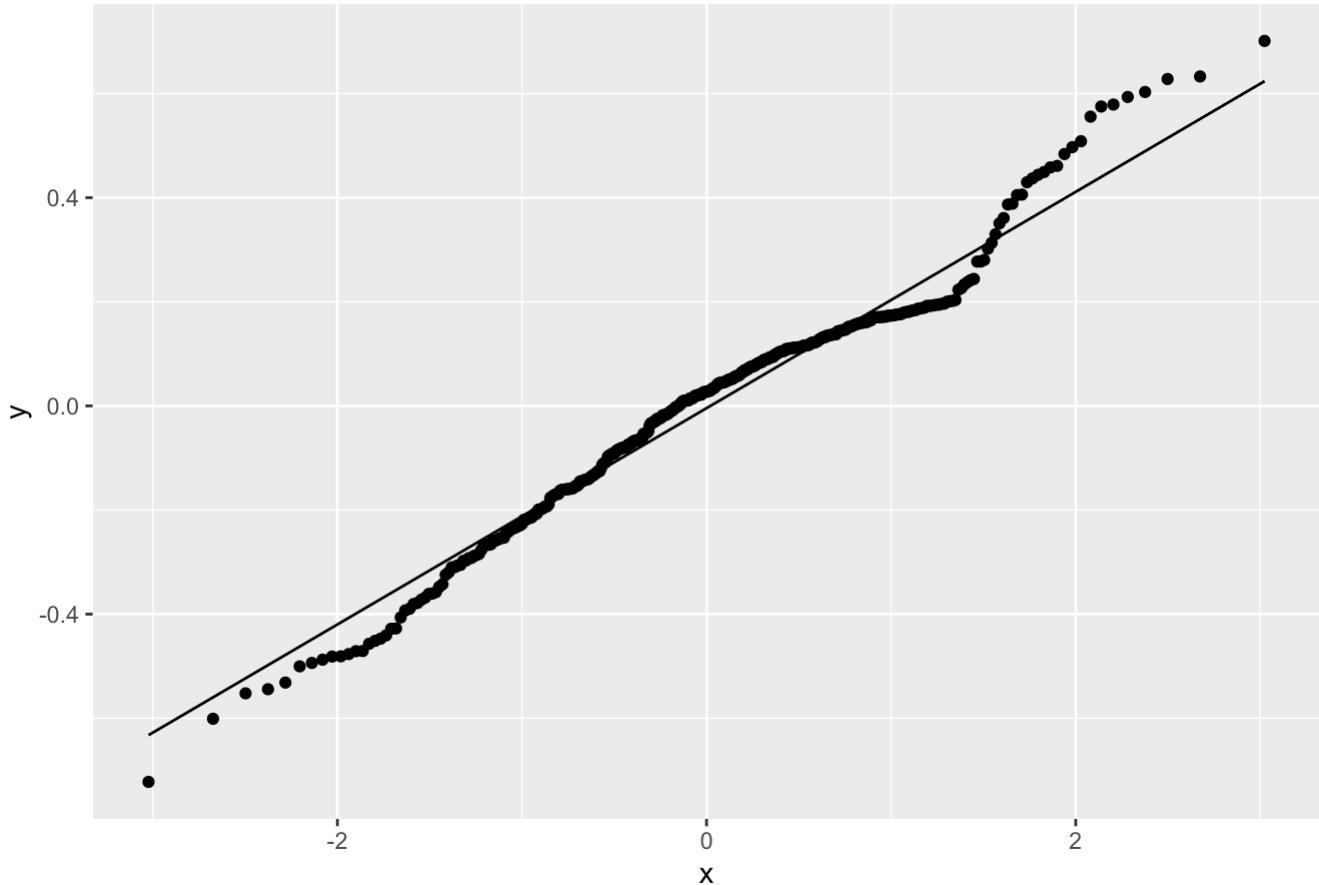
```
diag_df <- data.frame(  
  fitted = fitted(lm_fit),  
  resid  = resid(lm_fit),  
  std_resid = rstandard(lm_fit)  
)  
  
#1) Residuals vs Fitted  
ggplot(diag_df, aes(x = fitted, y = resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, linetype = "dashed") +  
  labs(title = "Residuals vs Fitted",  
       x = "Fitted values",  
       y = "Residuals")
```

Residuals vs Fitted



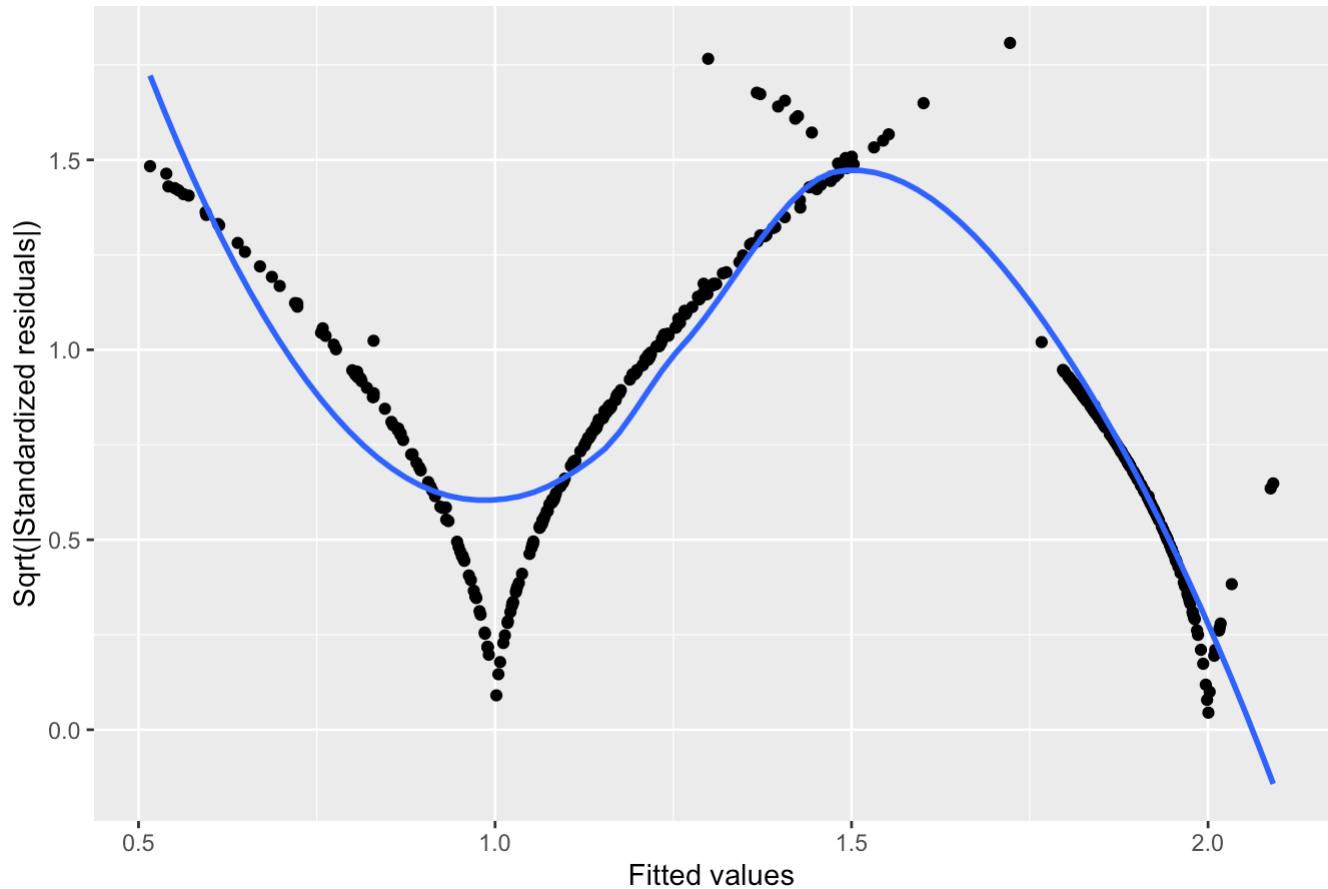
```
#2) QQ Plot  
ggplot(diag_df, aes(sample = resid)) +  
  stat_qq() +  
  stat_qq_line() +  
  labs(title = "Normal Q-Q Plot")
```

Normal Q-Q Plot



```
#3) Scale-Location (sqrt |resid|)  
ggplot(diag_df, aes(x = fitted, y = sqrt(abs(std_resid)))) +  
  geom_point() +  
  geom_smooth(se = FALSE) +  
  labs(title = "Scale-Location",  
       x = "Fitted values",  
       y = "Sqrt(|Standardized residuals|)")
```

Scale-Location



```
plot(lm_fit, which = 5)
```

